

Akshay K. Jagadish

Postdoctoral Research Fellow, Princeton AI Lab

📍 Princeton, USA | 📩 akshaykjagadish@gmail.com | 🏷 akjagadish.github.io | 🐾 akjagadish | 🐦 akjagadish | 🎓 Scholar

Education

Postdoctoral Research Fellow

Princeton, USA

2025 - ongoing

Princeton University

Natural and Artificial Minds (NAM) Initiative, Princeton AI Lab

Ph.D. in Computer Science

University of Tübingen · Max Planck Institute for Biological Cybernetics · Helmholtz Munich

Tübingen, Germany

2021 - 2025

Thesis: Meta-Learning: A Unifying Framework for Testing Theories of Human Learning

M.Sc. in Computational Neuroscience

Tübingen, Germany

2018 - 2020

University of Tübingen · Max Planck Institute for Biological Cybernetics

Thesis: Compositional Reinforcement Learning in Minds and Machines

B.Tech. in Electrical and Electronics Engineering

Surathkal, India

2013 - 2017

National Institute of Technology Karnataka · Ecole Polytechnique Federale de Lausanne

Thesis: Structural and Functional Correlates of Personality

Research Experience

Ph.D. Thesis

Tübingen, Germany

Apr. 2021 - Aug. 2025

University of Tübingen · Max Planck Institute for Biological Cybernetics · Helmholtz Munich

Computational Principles of Intelligence Lab · Advisors: Dr. Marcel Binz and Prof. Eric Schulz

Natural Minds

- Proposed a new framework for building computational cognitive models termed as “Meta-learned models of cognition” [PDF].
- Demonstrated that resource rational meta-learned inference explains zero-shot compositional reasoning in humans [PDF].
- Illustrated that ecologically rational meta-learned inference explains human category learning [PDF].
- Showed that humans and large language models display symmetric belief updating [PDF].
- Developed bounded ecologically rational meta-learned inference to explore the role of bounded resources and ecological priors in human learning [PDF].
- Demonstrated that human learning and decision making in 15 different experiments can be explained by meta-learning ecological priors from large language mdoels [PDF].

Artificial Minds

- Illustrated that Inducing anxiety in large language models increases exploration and bias [PDF].
- Demonstrated that narrative of traumatic experiences increases state anxiety in large language models, but using mindfulness-based techniques can help alleviate the same [PDF].
- Used sparse autoencoders to reveal temporal difference learning in large language mdules [PDF].

Automated scientific discovery

- Contributed towards building the first foundation model of human cognition [PDF].
- Generated cognitive models with large language models [PDF].
- Automated scientific minimization of regret to discover cognitive models [PDF].

Masters Thesis

Tübingen, Germany

Jun. 2020 - Dec. 2020

Max Planck Institute for Biological Cybernetics

Computational Principles of Intelligence Lab · Advisors: Dr. Marcel Binz and Prof. Eric Schulz

- Built meta-reinforcement learning agents that can discover and compose latent structures underlying rewards of structured bandit tasks, and further, generalize to novel structures unseen during training.

Graduate Research Assistant

Tübingen, Germany

Nov. 2018 - Mar. 2021

University of Tübingen

Sinz Lab · Advisors: Prof. Fabian Sinz and Prof. Edgar Walker

- Built a factor analysis model on top of a convolutional neural network, which predicts stimulus-based neural activity, to recover non-stimulus-related latent brain states [PDF].
- Developed a novel parameter-efficient readout, called a Gaussian readout, that maps nonlinear features learned by the deep convolutional network to the response of each neuron [PDF].

Graduate Research Assistant

Max Planck Institute for Biological Cybernetics

Computational Neuroscience Lab · Advisor: Prof. Peter Dayan

Tübingen, Germany

Nov. 2019 - Feb. 2020

- Conducted a literature review on the role of Dopamine in reward-based learning [PDF].

- Analyzed behavior (choices and reaction times) of monkeys, whose dopamine receptors were pharmacologically stimulated, during a rule-based categorization task [PDF].

AI Researcher

Wadhwani Institute for Artificial Intelligence

AI for Social Impact · Advisor: Dr. Rahul Panicker

Mumbai, India

May 2018 - Sep. 2018

- Developed a model based on deep-learning that predicts weight of an object from its images [NDA].

Postbaccalaureate Research Assistant

University of Minnesota, Twin-cities

Computational Visual Neuroscience Lab · Advisor: Prof. Kendrick Kay

Minnesota, USA

Jul. 2017 - Feb. 2018

- Developed a generative model that factors in the bottom-up, stimulus-driven, and top-down, goal-driven attentional state to characterize cortical fMRI responses for various stimuli and attentional loci combinations.

Undergraduate Thesis

Ecole Polytechnique Federale de Lausanne

Medical Image Processing Lab · Advisors: Prof. D. van de Ville and Prof. P. Giannakopoulos

Lausanne, Switzerland

Aug. 2016 - May 2017

- Investigated the relationship between structural and functional connectivity measures derived from MRI, and Neuroticism Extroversion Openness Personality Inventory-Revised (NEOPI) personality traits [PDF].

Undergraduate Research Assistant

Indian Institute of Science

Computational Tomography Lab · Advisor: Prof. Kasi Rajgopal

Bangalore, India

May 2015 - May 2017

- Developed an algorithm, called k-ABC, based on the artificial bee colony algorithm to come up with an optimal variable density sampling scheme for the compressed sensing-based reconstruction of Magnetic Resonance (MR) images [PDF].

Selected Publications

* equal contribution, # alphabetical ordering

Jagadish, A. K., Binz, M. & Schulz, E. (2025). Meta-learning ecological priors from large language models captures human learning and decision making. *Under Review*.

Rmus, M., **Jagadish, A. K.***, Mathony, M., & Schulz, E. (2025). Generating Computational Cognitive Models using Large Language Models. *In Advances in Neural Information Processing Systems (NeurIPS)* [PDF]

Binz, M., ..., **Jagadish, A. K.**#, ..., & Schulz, E. (2024). Centaur: a foundation model of human cognition. *Nature* [PDF]

Demirican, C.* , Saanum, T.* , **Jagadish, A. K.**, Binz, M., & Schulz, E. (2025). Sparse Autoencoders Reveal Temporal Difference Learning in Large Language Models. *In the 13th International Conference on Learning Representations (ICLR)* [PDF]

Jagadish, A. K., Thalmann, M., Coda-Forno, J., Schulz, E., & Binz, M. (2024). Human-like Category Learning by Injecting Ecological Priors from Large Language Models into Neural Networks. *Proceedings of the 41st International Conference on Machine Learning (ICML)* [PDF]

Schubert, J., **Jagadish, A. K.**, Binz, M., & Schulz, E. (2024). In-context learning agents are asymmetric belief updaters. *Proceedings of the 41st International Conference on Machine Learning (ICML)* [PDF]

Jagadish, A. K., Binz, M., Saanum, T., Wang, J. X., & Schulz, E. (2024). Zero-shot compositional reasoning in a reinforcement learning setting. [PDF]

Ben-Zion, Z., Witte, K.* , **Jagadish, A. K.***, Duek, O., Harpaz-Rotem, I., Khorsandian, M., ... & Spiller, T. R. (2024). “Chat-GPT on the Couch”: Assessing and Alleviating State Anxiety in Large Language Models. *npj Digital Medicine* [PDF]

Coda-Forno, J.* , Witte, K.* , **Jagadish, A. K.**, Binz, M., Akata, Z., & Schulz, E. (2023). Inducing anxiety in large language models increases exploration and bias. [PDF];

Binz, M., Dasgupta, I., **Jagadish, A. K.**, Botvinick, M., Wang, J.X., & Schulz, E. (2023). Meta-learned models of cognition. *Behavioral and Brain Sciences*, 1-38. [PDF]

Bashiri, M.* , Walker, E.* , Lurz, K. K., **Jagadish, A. K.**, Muhammad, T., Ding, Z., ... & Sinz, F. (2021). A flow-based latent state generative model of neural population responses to natural images. *In Advances in Neural Information Processing Systems (NeurIPS), 34, 15801-15815.* [PDF]

Lurz, K. K., Bashiri, M., Willeke, K. F., **Jagadish, A. K.**, Wang, E., Walker, E. Y., ... & Sinz, F. (2021). Generalization in data-driven models of primary visual cortex. *In International Conference on Learning Representations (ICLR)* [PDF]

Honors, Awards, & Fellowships

2025	Top Reviewer , Top 10 % reviewer (main track) for Neural Information Processing Systems (NeurIPS)	USA
2025	Ph.D. with Magna cum Laude , for exceptional research conducted during doctoral studies	Germany
2024	Top Reviewer , Top 10 % reviewer (main track) for Neural Information Processing Systems (NeurIPS)	Canada
2024	ELLIS Winter School on Foundation Models , Top 40 students selected from 607 to present their research	Netherlands
2023	Analytical Connectionism , Top 35 students selected to attend the summer course at the Gatsby Computational Neuroscience Unit in University College London	UK
2020	SMARTSTART Fellowship in Computational Neuroscience , Top 15 students awarded a travel budget of 1000 euros and mentorship from Prof. Peter Dayan and Prof. Fabian Sinz	Germany
2019	Dean's List , Top 3 performers in the summer semester for M.Sc. program in Computational Neuroscience	Germany
2018	Dean's List , Top 3 performers in the winter semester for M.Sc. program in Computational Neuroscience	Germany
2018	Max Planck Society Scholarship , Top 5 students selected to undertake M.Sc in Computational Neuroscience at the University of Tübingen	Germany
2017	Harvard Young Scientist Development Program , Top 25 students selected for training in neuroscience	USA India
2016	Summer Research Program , Top 20 students funded to conduct research at EPFL	Switzerland
2016	Summer Research Fellowship Program , Top 10 % students funded to conduct research at the Indian Institute of Science	India
2013	Ranked Top 0.1% , Karnataka Common Entrance Test among 150,000 students	India
2013	Ranked Top 0.1% , COMED-K among 200,000 students	India
2011	Most Consistent Performer of the Batch , High school at Presidency School	India

Press

2025	New York Times , A science-education piece called "Digital Therapists Get Stressed Too, Study Finds" on our article "Chat-GPT on the Couch": Assessing and Alleviating State Anxiety in Large Language Models [URL]	AoE
2025	Fortune Magazine , Interview to discuss our article "Chat-GPT on the Couch": Assessing and Alleviating State Anxiety in Large Language Models [URL]	AoE
2025	ScienceDaily , Accessible take on our article "Chat-GPT on the Couch": Assessing and Alleviating State Anxiety in Large Language Models [URL]	AoE

Invited talks

2025	Princeton University , Natural and Artificial Minds (NAM) talk series in the the Princeton AI Lab	USA
2025	Princeton University , Workshop on Automated Discovery of Mind and Brain	USA
2025	University of Osnabrück , Lab meeting at the Laboratory for Automated Scientific Discovery of Mind and Brain	Germany
2024	Annual Meeting of the Cognitive Science Society (CogSci) , Workshop on "In-context learning in natural and artificial intelligence"	Netherlands
2024	Annual Meeting of the Cognitive Science Society (CogSci) , Workshop on "Compositionality in minds, brains and machines: A unifying goal that cuts across cognitive sciences"	Netherlands
2024	Indian Institute of Science , Seminar talk at the Center for Neuroscience	India
2024	Princeton University , Lab meeting at Computational Cognitive Science Lab	USA
2023	Helmholtz München , Joint lab retreat with Explainable Machine Learning Lab	Germany
2023	University of Oxford , Lab meeting at the Human Information Processing Lab	United Kingdom
2022	Max Planck Institute for Human Cognitive and Brain Sciences , Joint lab retreat with Doeller Lab	Germany
2021	Stanford University , Joint lab retreat with Human Information Processing and Causality in Cognition Lab	USA
2019	University of Tübingen , Workshop on "Causality in Neuroscience" at Neuroscience Conference for Young Scientists	Germany

Organization

2024	Co-organizer , “Connecting Minds and Machines a Foundation Model Approach to Learning” symposium at Helmholtz Pioneer Campus, Munich	Germany
2024	Co-organizer , “In-context learning in natural and artificial intelligence” workshop at CogSci-2024	Netherlands
2023	Co-organizer , Laboratory Retreat of Computational Principles of Intelligence Lab	Germany
2022	Co-organizer , Computation and Cognitive Tübingen Summer School (CaCTüs) aimed specifically at young scientists held back by personal, financial, regional or societal constraints.	Germany
2019	Co-organizer , “Causality in Neuroscience” Workshop at Neuroscience Conference for Young Scientists	Germany
2020	Volunteer , Machine Learning Summer School (MLSS) held at MPI for Intelligent Systems, Tübingen	Germany

Supervision

2025	Doctoral Student , Liyi Zhang’s doctoral research project on “Distilling inductive reasoning into LLMs” at the Princeton University	USA
2025	Doctoral Student , Solim Legris’s doctoral research project on “Neuro-symbolic models for ARC benchmark” at the Princeton University	USA
2025	Doctoral Student , Younes Strittmatter’s doctoral research project on “Tools for automated scientific discovery” at the Princeton University	USA
2025	Master Student , Daniel Braga’s master research project on “Automated representational structure discovery” at the Princeton University	USA
2024	Master Student , Elif Kara’s master research project on “Human decision-making in the wild” at the University of Münich	Germany
2023	Master Student , Johannes Schubert’s master thesis on “Rational Analysis of Optimism Bias” at the University of Tübingen. Converted to a publication at the <i>International Conference on Machine Learning (ICML)</i> [PDF]	Germany

Teaching

2023	Teaching Assistant , “Computational Cognitive Science” course at the Graduate Training Center for Neuroscience, University of Tübingen	Germany
2022	Tutor , Tutorial on “Meta-Reinforcement Learning” at the Max Planck Institute for Biological Cybernetics	Germany

Service

2025-	Reviewer , Transactions on Machine Learning Research (TMLR)	AoE
2024-	Reviewer , Neural Information Processing Systems (NeurIPS)	AoE
2023-	Reviewer , International Conference on Learning Representations (ICLR)	AoE
2025-	Reviewer , International Conference on Machine Learning (ICML)	AoE
2022-	Reviewer , Annual Meeting of the Cognitive Science Society (CogSci)	AoE
2023-	Reviewer , Cognitive Computational Neuroscience (CCN)	AoE